

SOLVING PROBLEMS BY SEARCHING

CHAPTER 3

In which we see how an agent can find a sequence of actions that achieves its goals when no single action will do.

Outline

- ◇ Problem-solving agents
- ◇ Problem formulation
- ◇ Basic search algorithms
- ◇ Informed (heuristic) search strategies

Example: Romania

On holiday in Romania; currently in Arad.

Flight leaves tomorrow from Bucharest

Formulate goal:

be in Bucharest

Formulate problem: (granularity level)


states: various cities

actions: drive between cities

Find solution:

sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

Example: Romania



Properties of the environment

observable - so the agent always knows the current state.

discrete - so at any given state there are only finitely many actions to choose from.

known, so the agent knows which states are reached by each action.

deterministic - so each action has exactly one outcome

Problem types

Deterministic, fully observable \implies single-state problem

Agent knows exactly which state it will be in; solution is a sequence

Non-observable \implies conformant problem

Agent may have no idea where it is; solution (if any) is a sequence

Nondeterministic and/or partially observable \implies contingency problem

percepts provide **new** information about current state

solution is a **contingent plan** or a **policy**


often **interleave** search, execution

The agent might plan to drive from Arad to Sibiu and then to Rimnicu Vilcea but may also need to have a contingency plan in case it arrives by accident in Zerind instead of Sibiu.

Unknown state space \implies exploration problem (“online”)

Example: vacuum world

Single-state, start in #5. Solution??
Conformant




Example: vacuum world

Single-state, start in #5. Solution??

[*Right, Suck*]

Conformant, start in $\{1, 2, 3, 4, 5, 6, 7, 8\}$

e.g., *Right* goes to $\{2, 4, 6, 8\}$. Solution??



Example: vacuum world

Single-state, start in #5. Solution??

[*Right, Suck*]

Conformant, start in {1, 2, 3, 4, 5, 6, 7, 8}

e.g., *Right* goes to {2, 4, 6, 8}. Solution??


[*Right, Suck, Left, Suck*]

Contingency, start in #5

Murphy's Law: *Suck* can dirty a clean carpet

Local sensing: dirt, location only.

Solution??



Example: vacuum world

Single-state, start in #5. Solution??

[*Right, Suck*]

Conformant, start in {1, 2, 3, 4, 5, 6, 7, 8}

e.g., *Right* goes to {2, 4, 6, 8}. Solution??

[*Right, Suck, Left, Suck*]


Contingency, start in #5

Murphy's Law: *Suck* can dirty a clean carpet

Local sensing: dirt, location only.

Solution??

[*Right, if dirt then Suck*]



Single-state problem formulation

◇ A **problem** is defined by five items:

1. **initial state** e.g., “at Arad”
2. **actions**: $\text{Actions}(s)$ returns applicable actions in s : ($\text{Go}(\text{Sibiu})$, $\text{Go}(\text{Timisoara})$, $\text{Go}(\text{Zerind})$)
3. **transition model** - set of action–state pairs (successor function $\text{Result}(s,a)$):
e.g., $\text{Result}(\text{In}(\text{Arad}), \text{Go}(\text{Zerind})) = \text{In}(\text{Zerind})$
4. **goal test** - determines if whether a given state is a goal state
explicit, e.g., $xs = \text{“at Bucharest”}$
implicit, e.g., $xs = \text{checkmate}$
5. **path cost** (additive) - reflects agent’s own performance measure
e.g., sum of distances, number of actions executed, etc.
 $c(s, a, s')$ is the **step cost**, assumed to be ≥ 0

◇ A **solution** is a sequence of actions leading from the initial state to a goal state

Selecting a state space

Real world is absurdly complex

⇒ state space must be **abstracted** for problem solving

(Abstract) state = set of real states

(Abstract) action = complex combination of real actions

e.g., “Arad → Zerind” represents a complex set
of possible routes, detours, rest stops, etc.


For guaranteed realizability, **any** real state “in Arad”
must get to **some** real state “in Zerind”

(Abstract) solution =

set of real paths that are solutions in the real world

Each abstract action should be “easier” than the original problem!

Example: vacuum world state space graph




states??

actions??

goal test??

path cost??

Example: vacuum world state space graph




states??: integer dirt and robot locations (ignore dirt **amounts** etc.)

actions??

goal test??

path cost??

Example: vacuum world state space graph



states??: integer dirt and robot locations (ignore dirt **amounts** etc.)


actions??: *Left, Right, Suck, NoOp*

transition model??:

goal test??

path cost??

Example: vacuum world state space graph



states??: integer dirt and robot locations (ignore dirt **amounts** etc.)


actions??: *Left, Right, Suck, NoOp*

transition model??:

goal test??: no dirt

path cost??

Example: vacuum world state space graph



states??: integer dirt and robot locations (ignore dirt **amounts** etc.)

actions??: *Left*, *Right*, *Suck*, *NoOp*

transition model??: Fig.

goal test??: no dirt

path cost??: 1 per action (0 for *NoOp*)

Vacuum world vs. real world:

Example: The 8-puzzle

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| 1 | 2 | 3 |
| 4 | 5 | 6 |
| 7 | 8 | |

Goal State

states??

actions??

transition model??

goal test??

path cost??

Example: The 8-puzzle

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| 1 | 2 | 3 |
| 4 | 5 | 6 |
| 7 | 8 | |

Goal State

states??: integer locations of tiles (ignore intermediate positions)

actions??

transition model??

goal test??

path cost??

Example: The 8-puzzle

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| 1 | 2 | 3 |
| 4 | 5 | 6 |
| 7 | 8 | |

Goal State

states??: integer locations of tiles (ignore intermediate positions)

actions??: move blank left, right, up, down (ignore unjamming etc.)

goal test??

path cost??

Example: The 8-puzzle

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| 1 | 2 | 3 |
| 4 | 5 | 6 |
| 7 | 8 | |

Goal State

states??: integer locations of tiles (ignore intermediate positions)

actions??: move blank left, right, up, down (ignore unjamming etc.)

goal test??: = goal state (given)

path cost??

Example: The 8-puzzle

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| 1 | 2 | 3 |
| 4 | 5 | 6 |
| 7 | 8 | |

Goal State


states??: integer locations of tiles (ignore intermediate positions)

actions??: move blank left, right, up, down (ignore unjamming etc.)

goal test??: = goal state (given)

path cost??: 1 per move

Example: robotic assembly



states??: real-valued coordinates of robot joint angles
parts of the object to be assembled

actions??: continuous motions of robot joints

goal test??: complete assembly

path cost??: time to execute

Outline

- ◇ Problem-solving agents
- ◇ Problem formulation
- ◇ Uninformed search algorithms
- ◇ Informed (heuristic) search strategies

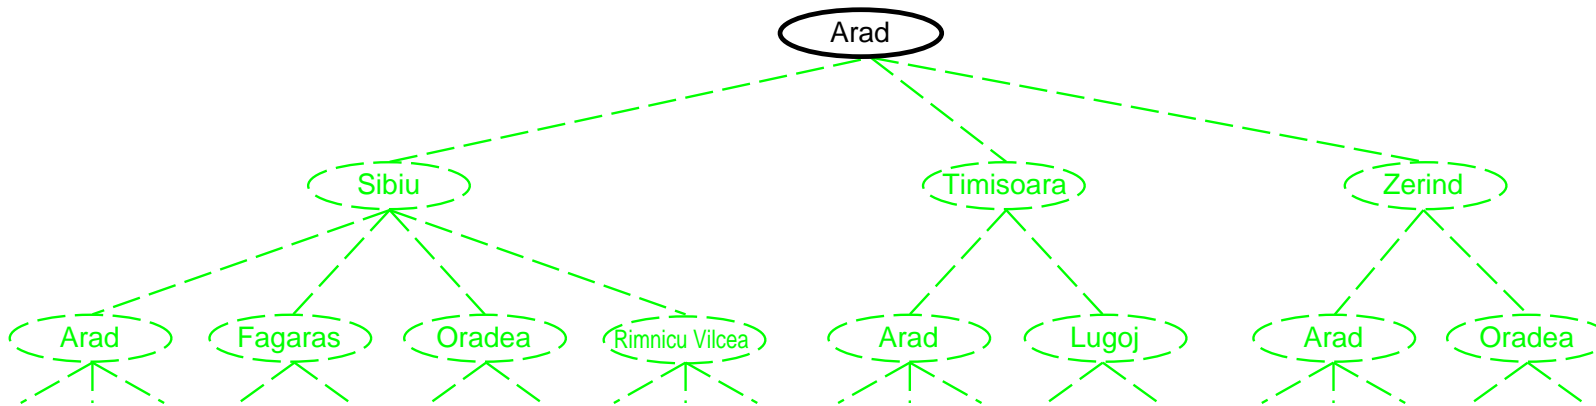
Tree search algorithms

Basic idea:

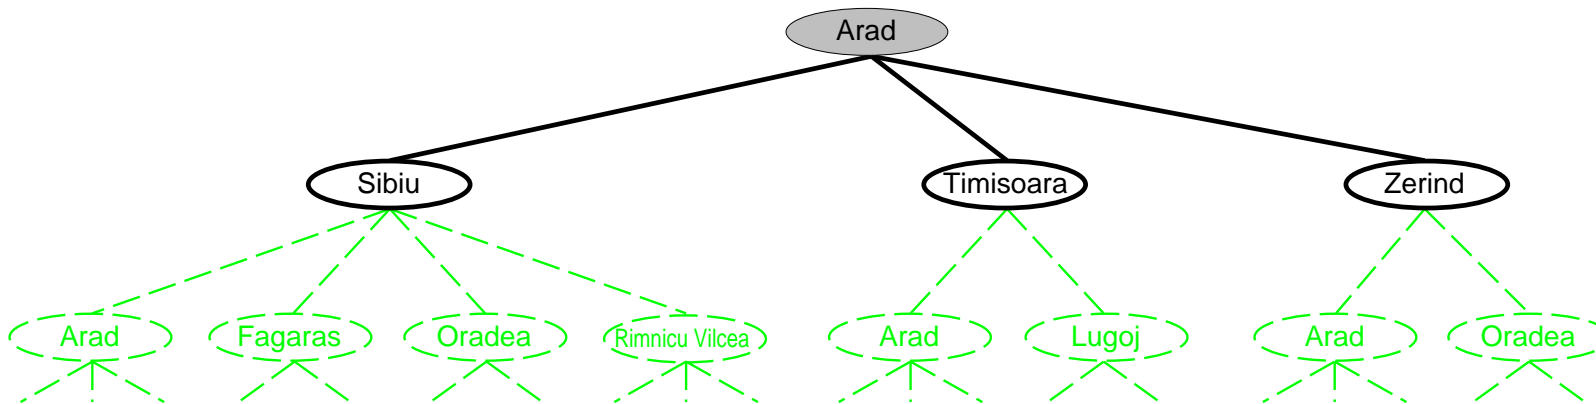
offline, simulated exploration of state space
by generating successors of already-explored states
(a.k.a. **expanding** states)

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

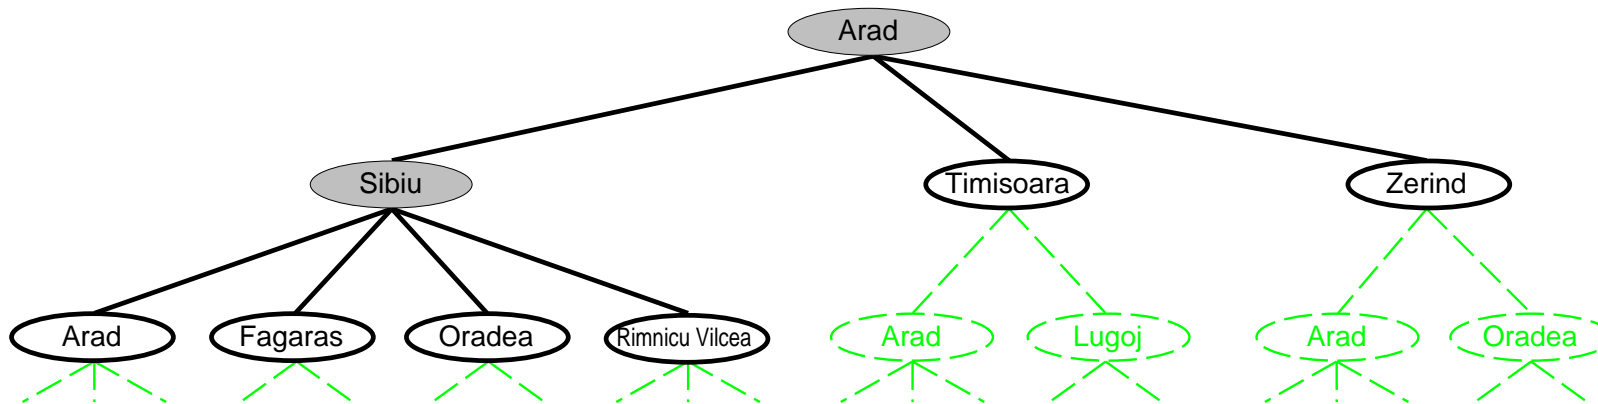
Tree search example



Tree search example



Tree search example




Implementation: states vs. nodes

A **state** is a (representation of) a physical configuration

A **node** is a data structure constituting part of a search tree

includes **parent**, **children**, **depth**, **path cost** $g(x)$

States do not have parents, children, depth, or path cost!



The EXPAND function creates new nodes, filling in the various fields and using the SUCCESSORFN of the problem to create the corresponding states.

Two different nodes can contain the same world state

Implementation: general tree search

```

function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE(node)) then return node
    fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)

```

```

function EXPAND(node, problem) returns a set of nodes
  successors  $\leftarrow$  the empty set
  for each action, result in SUCCESSOR-FN(problem, STATE[node]) do
    s  $\leftarrow$  a new NODE
    PARENT-NODE[s]  $\leftarrow$  node; ACTION[s]  $\leftarrow$  action; STATE[s]  $\leftarrow$  result
    PATH-COST[s]  $\leftarrow$  PATH-COST[node] + STEP-COST(STATE[node], action,
    result)
    DEPTH[s]  $\leftarrow$  DEPTH[node] + 1
    add s to successors
  return successors

```

Search strategies

A strategy is defined by picking the **order of node expansion**

Strategies are evaluated along the following dimensions:

completeness—does it always find a solution if one exists?

time complexity—number of nodes generated/expanded

space complexity—maximum number of nodes in memory

optimality—does it always find a least-cost solution?

Time and space complexity are measured in terms of

b —maximum branching factor of the search tree

d —depth of the least-cost solution (root-goal)

m —maximum depth of any path in the state space (may be ∞)

search cost: the amount of time taken by search

solution cost: the total length of the path

total cost = search cost + solution cost

Uninformed search strategies

Uninformed strategies use only the information available in the problem definition

Breadth-first search

Uniform-cost search

Depth-first search

Depth-limited search


Iterative deepening search

Breadth-first search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end




Breadth-first search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end




Breadth-first search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end




Breadth-first search

Expand shallowest unexpanded node

Implementation:

fringe is a FIFO queue, i.e., new successors go at end



Breadth-first search on a graph

function BREADTH-FIRST-SEARCH(*problem*) **returns** a solution, or failure

node \leftarrow a node with STATE = *problem*.INITIAL-STATE, PATH-COST = 0

if *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)

frontier \leftarrow a FIFO queue with *node* as the only element

explored \leftarrow an empty set

loop do

if EMPTY?(*frontier*) **then return** failure

node \leftarrow POP(*frontier*) /* chooses the shallowest node in *frontier* */

 add *node*.STATE to *explored*

for each *action* **in** *problem*.ACTIONS(*node*.STATE) **do**

child \leftarrow CHILD-NODE(*problem*, *node*, *action*)

if *child*.STATE is not in *explored* or *frontier* **then**

if *problem*.GOAL-TEST(*child*.STATE) **then return** SOLUTION(*child*)

frontier \leftarrow INSERT(*child*, *frontier*)

function CHILD-NODE(*problem*, *parent*, *action*) **returns** a node

return a node with

 STATE = *problem*.RESULT(*parent*.STATE, *action*),

 PARENT = *parent*, ACTION = *action*,

 PATH-COST = *parent*.PATH-COST + *problem*.STEP-COST(*parent*.STATE, *action*)

Properties of breadth-first search

Complete??

Properties of breadth-first search

Complete?? Yes (if b is finite)

Time??

Properties of breadth-first search

Complete?? Yes (if b is finite)

Time?? $1 + b + b^2 + b^3 + \dots + b^d = O(b^d)$, i.e., exp. in d

Space??

Properties of breadth-first search

Complete?? Yes (if b is finite)

Time?? $1 + b + b^2 + b^3 + \dots + b^d = O(b^d)$, i.e., exp. in d

Space?? $O(b^d)$ (keeps every node from the frontier $O(b^d)$ and explored set $O(b^{d-1})$ in memory)

Optimal??

Properties of breadth-first search

Complete?? Yes (if b is finite)

Time?? $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^d)$, i.e., exp. in d

Space?? $O(b^{d+1})$ (keeps every node in memory)


Optimal?? Yes (if cost = 1 per step); not optimal in general

| Depth | Nodes | Time | Memory |
|-------|-----------|------------------|----------------|
| 2 | 110 | .11 milliseconds | 107 kilobytes |
| 4 | 11,110 | 11 milliseconds | 10.6 megabytes |
| 6 | 10^6 | 1.1 seconds | 1 gigabyte |
| 8 | 10^8 | 2 minutes | 103 gigabytes |
| 10 | 10^{10} | 3 hours | 10 terabytes |
| 12 | 10^{12} | 13 days | 1 petabyte |
| 14 | 10^{14} | 3.5 years | 99 petabytes |
| 16 | 10^{16} | 350 years | 10 exabytes |

Figure 3.13 Time and memory requirements for breadth-first search. The numbers shown assume branching factor $b = 10$; 1 million nodes/second; 1000 bytes/node.

Uniform-cost search

Expand the node with the lowest path cost $g(n)$



Implementation:

fringe = queue ordered by path cost, lowest first


Equivalent to breadth-first if step costs all equal

Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

Implementation:

fringe = LIFO queue, i.e., put successors at front




Depth-first search

Expand deepest unexpanded node

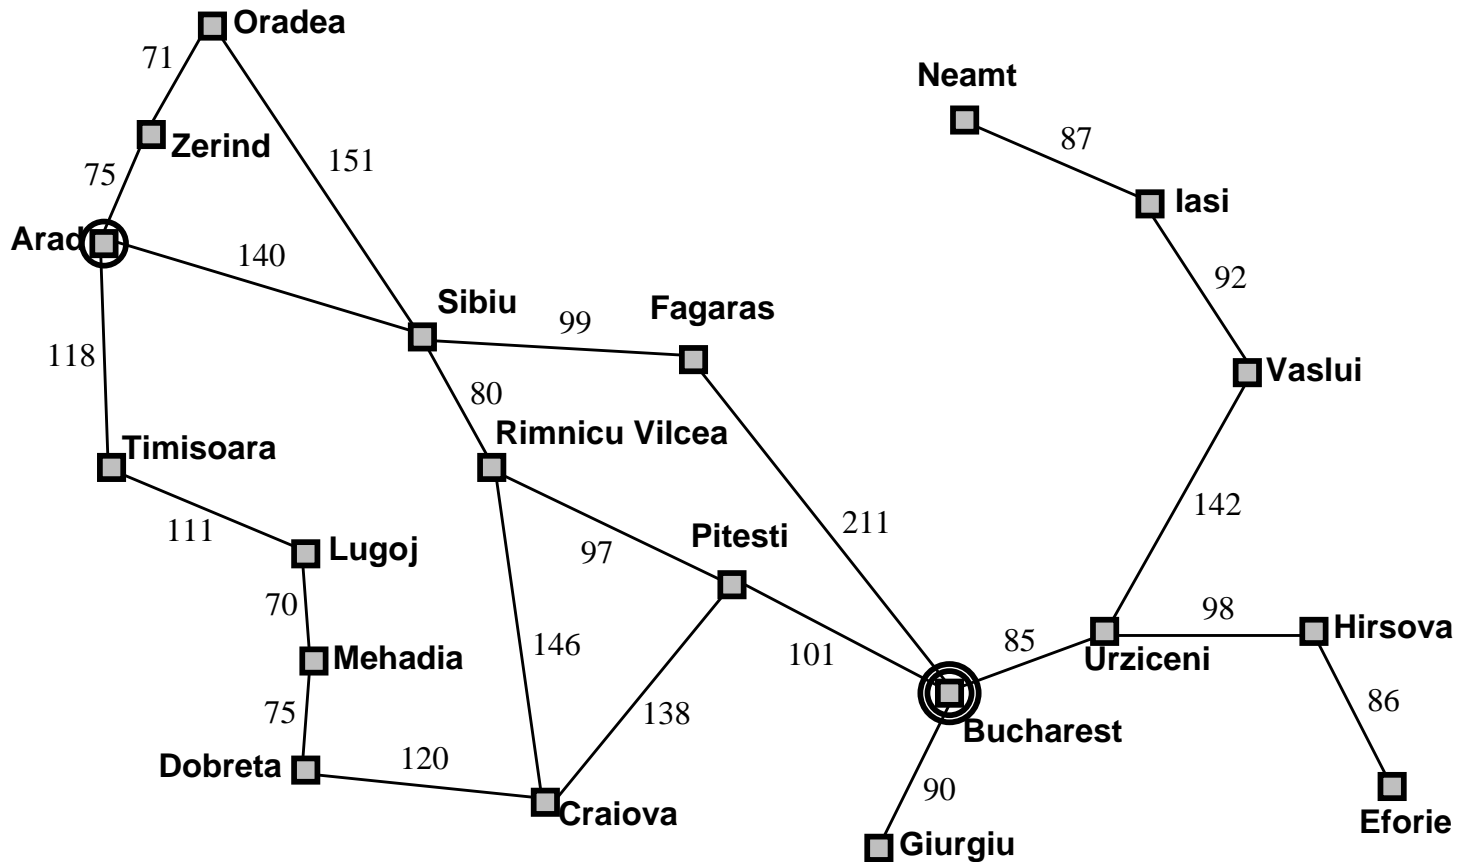
Implementation:

fringe = LIFO queue, i.e., put successors at front



Properties of depth-first search

Complete??



Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops
Arad-Sibiu-Arad-Sibiu....

Time??

Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path

⇒ complete in finite spaces

Time?? $O(b^m)$: terrible if m is much larger than d
but if solutions are dense, may be much faster than breadth-first

Space??

Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path


⇒ complete in finite spaces

Time?? $O(b^m)$: terrible if m is much larger than d

but if solutions are dense, may be much faster than breadth-first

Space?? $O(bm)$, i.e., linear space!

Optimal??



Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path


⇒ complete in finite spaces

Time?? $O(b^m)$: terrible if m is much larger than d

but if solutions are dense, may be much faster than breadth-first

Space?? $O(bm)$, i.e., linear space!

Optimal?? No (if J and C are goal nodes, the algorithm returns J)



Depth-limited search

= depth-first search with depth limit l , i.e., nodes at depth l have no successors

a new source of incompleteness ($l < d$) and nonoptimality ($l > d$)

Recursive implementation:

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff
    RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
    cutoff-occurred?  $\leftarrow$  false
    if GOAL-TEST(problem, STATE[node]) then return node
    else if DEPTH[node] = limit then return cutoff
    else for each successor in EXPAND(node, problem) do
        result  $\leftarrow$  RECURSIVE-DLS(successor, problem, limit)
        if result = cutoff then cutoff-occurred?  $\leftarrow$  true
        else if result  $\neq$  failure then return result
    if cutoff-occurred? then return cutoff else return failure
```

Iterative deepening search

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution
  inputs: problem, a problem
  for depth  $\leftarrow$  0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
end
```

Iterative deepening search $l = 0$

Limit = 0



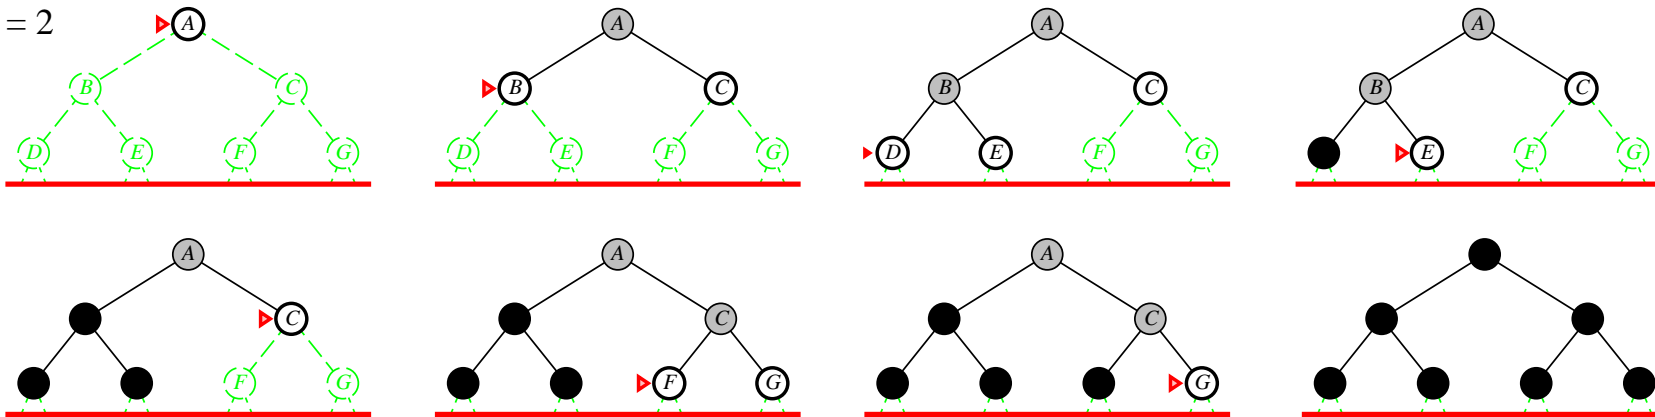
Iterative deepening search $l = 1$

Limit = 1




Iterative deepening search $l = 2$

Limit = 2



Iterative deepening search $l = 3$

Limit = 3



Properties of iterative deepening search

Complete??

Properties of iterative deepening search

Complete?? Yes

Time??

Properties of iterative deepening search

Complete?? Yes

Time?? - nodes generated $N(IDS)(d)b^1 + (d-1)b^2 + \dots + (1)b^d$
recall $N(BFS) = b^1 + b^2 + \dots + b^d$
b=10, d=4

Space??

Properties of iterative deepening search

Complete?? Yes

Time?? $(d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + b^d = O(b^d)$

Space?? $O(bd)$

Optimal??

Properties of iterative deepening search

Complete?? Yes

Time?? $(d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + b^d = O(b^d)$

Space?? $O(bd)$

Optimal?? Yes, if step cost = 1


Can be modified to explore uniform-cost tree

Numerical comparison for $b = 10$ and $d = 5$, solution at far right leaf:

$$N(\text{IDS}) = 50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$$

$$N(\text{BFS}) = 10 + 100 + 1,000 + 10,000 + 100,000 + 999,990 = 1,111,100$$

Bidirectional search



$$b^{d/2} + b^{d/2} < b^d$$

Replacing the goal test: check whether the frontiers intersect

Solution is not optimal

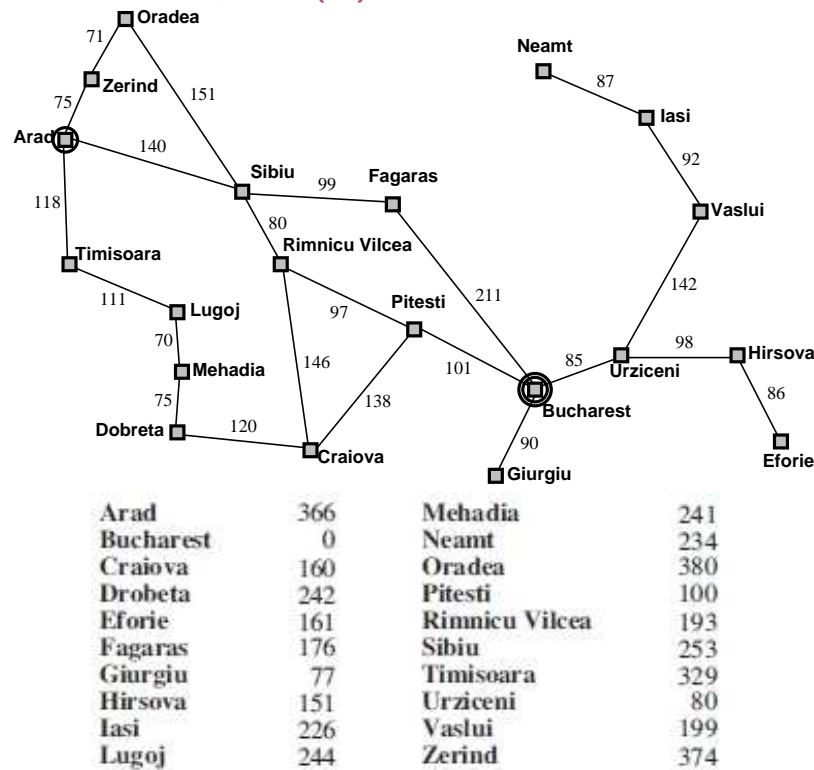
What if the goal is an abstract description: "no queen attacks another queen"?

Outline

- ◇ Problem-solving agents
- ◇ Problem formulation
- ◇ Basic search algorithms
- ◇ Informed (heuristic) search strategies

Greedy best first search

Heuristic function $h(n)$ = estimated cost of the cheapest path from the state at node n to a goal state, e.g., $h(n)$ = straight line distance




h_{sl_d} cannot be computed from the problem description itself

Search cost is minimal: without expanding a node not on the solution path. Is it optimal? Is it complete? (Iasi-Fagaras)


(a) The initial state



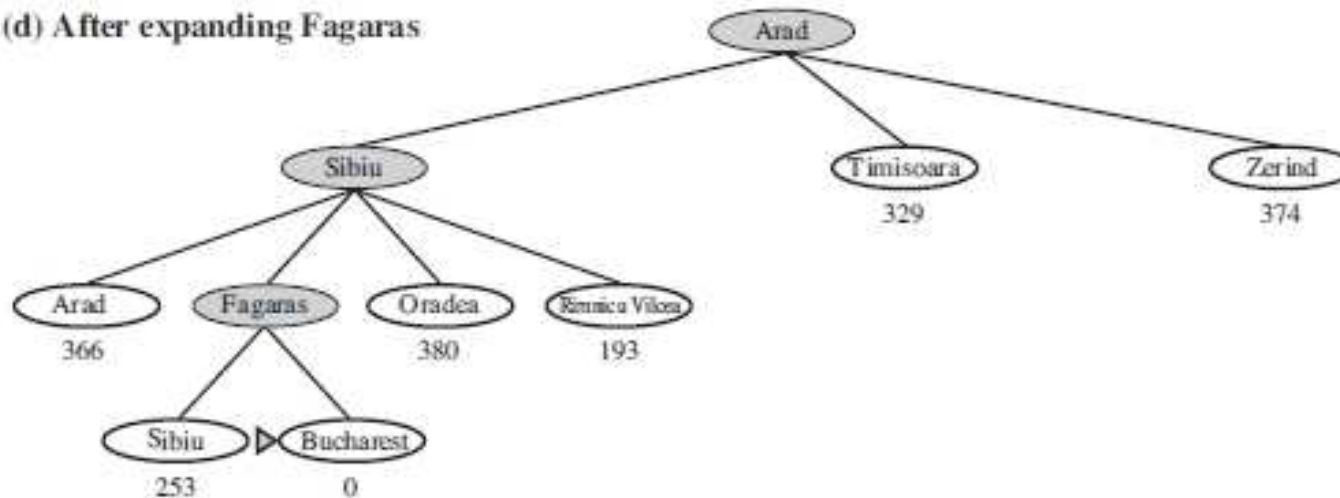
(b) After expanding Arad



(c) After expanding Sibiu



(d) After expanding Fagaras

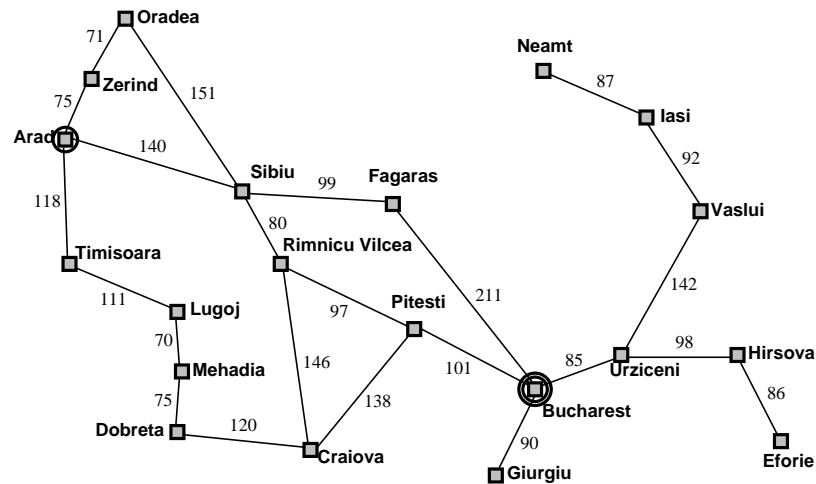


A* search: minimising the total estimated solution cost

$g(n)$ - the cost to reach the solution

$h(n)$ - the cost to get from the node to the goal

$f(n) = g(n) + h(n)$




| | | | |
|-----------|-----|----------------|-----|
| Arad | 366 | Mehadia | 241 |
| Bucharest | 0 | Neamt | 234 |
| Craiova | 160 | Oradea | 380 |
| Dobreta | 242 | Pitesti | 100 |
| Eforie | 161 | Rimnicu Vilcea | 193 |
| Fagaras | 176 | Sibiu | 253 |
| Giurgiu | 77 | Timisoara | 329 |
| Hirsova | 151 | Urziceni | 80 |
| Iasi | 226 | Vaslui | 199 |
| Lugoj | 244 | Zerind | 374 |

A* search: minimising the total estimated solution cost


(a) The initial state



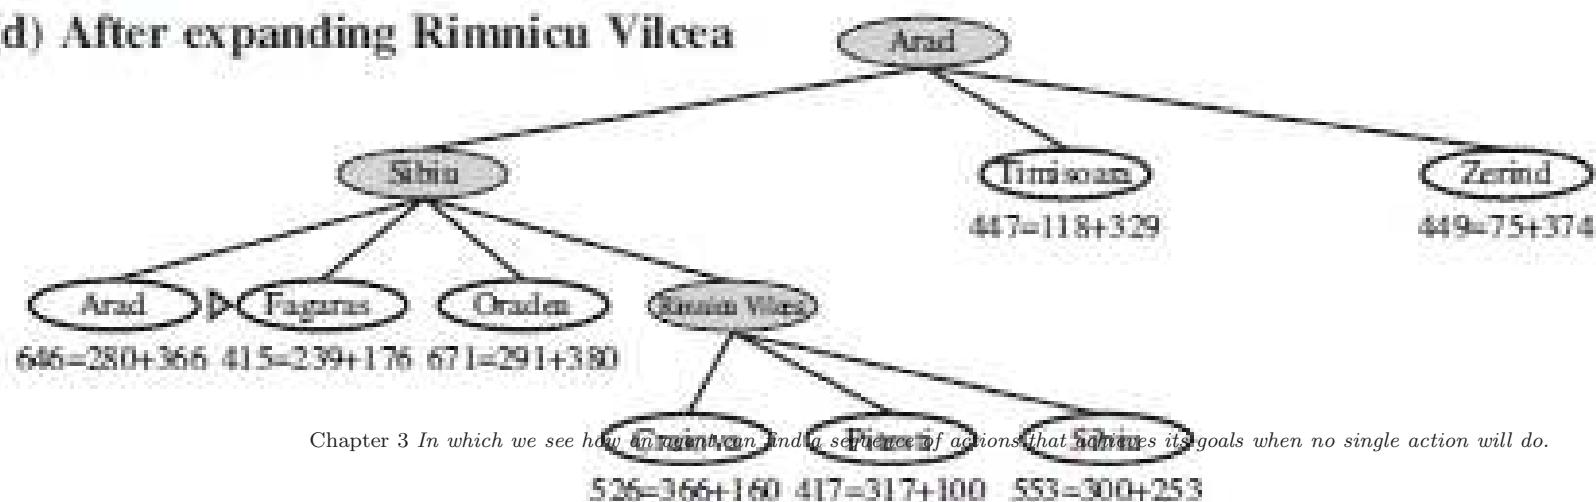
(b) After expanding Arad



(c) After expanding Sibiu




(d) After expanding Rimnicu Vilcea

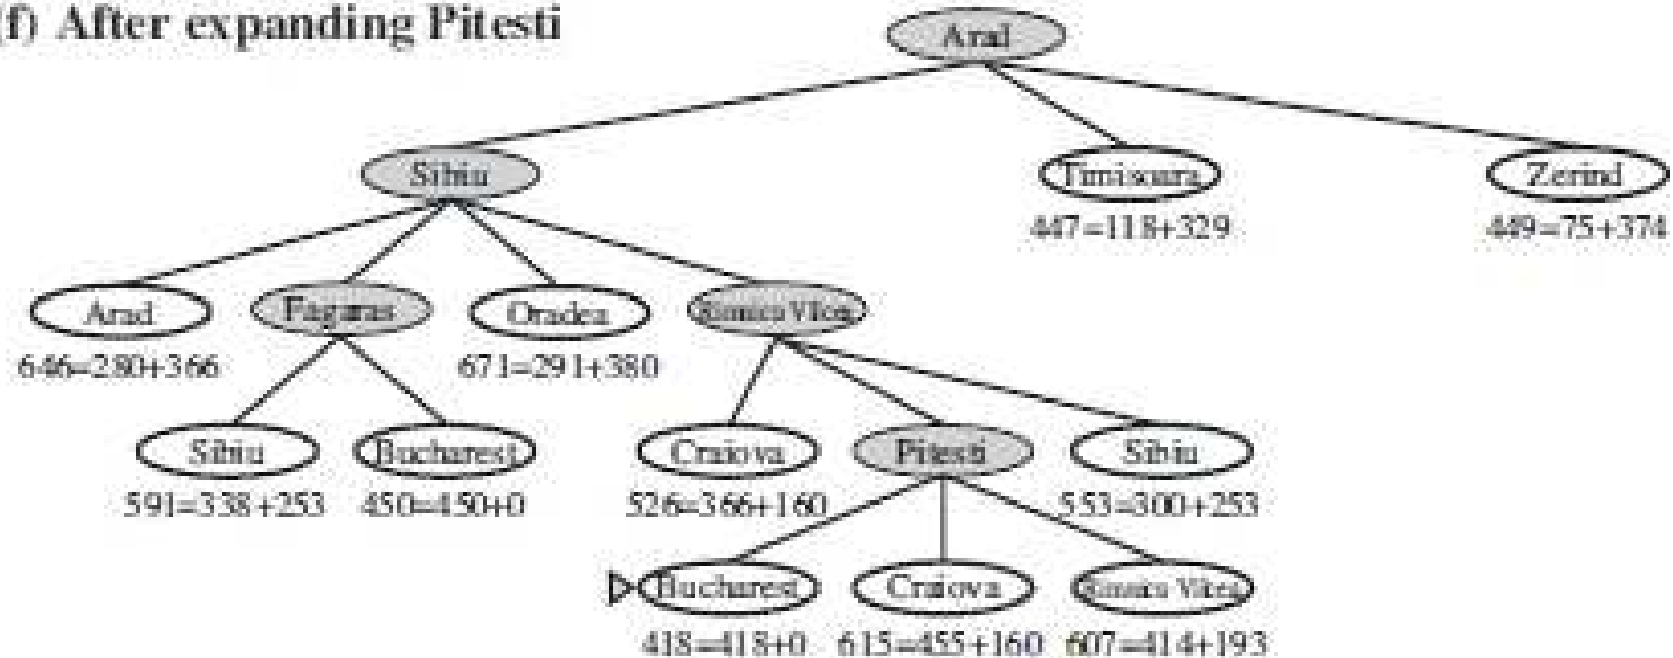


A* search: minimising the total estimated solution cost

(e) After expanding Fagaras

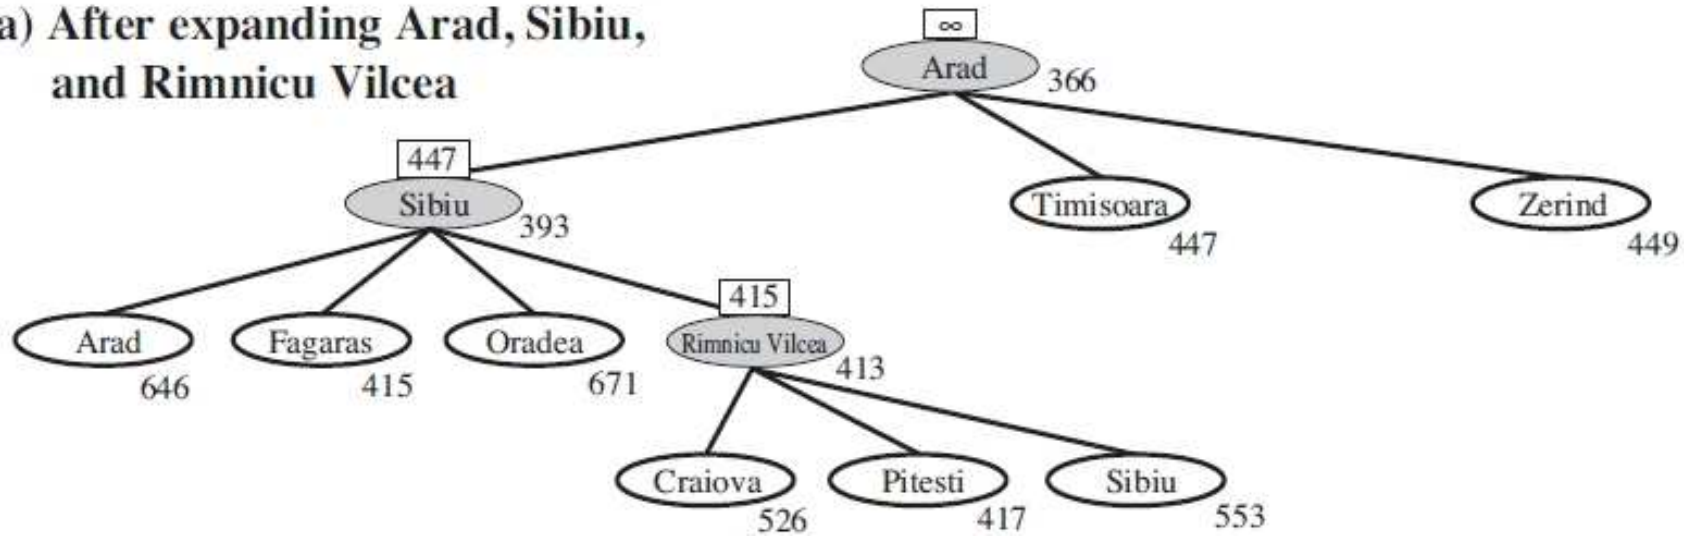


(f) After expanding Pitesti



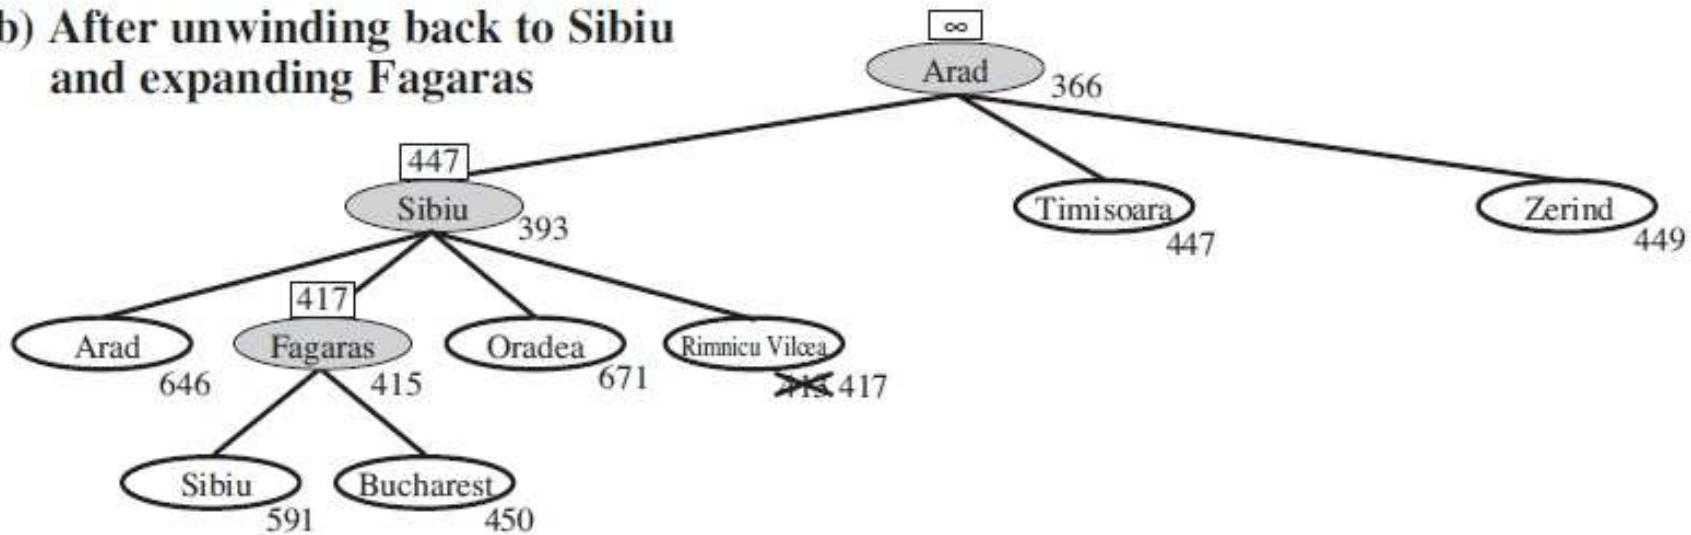
Memory-bounded heuristic search 1

(a) After expanding Arad, Sibiu, and Rimnicu Vilcea



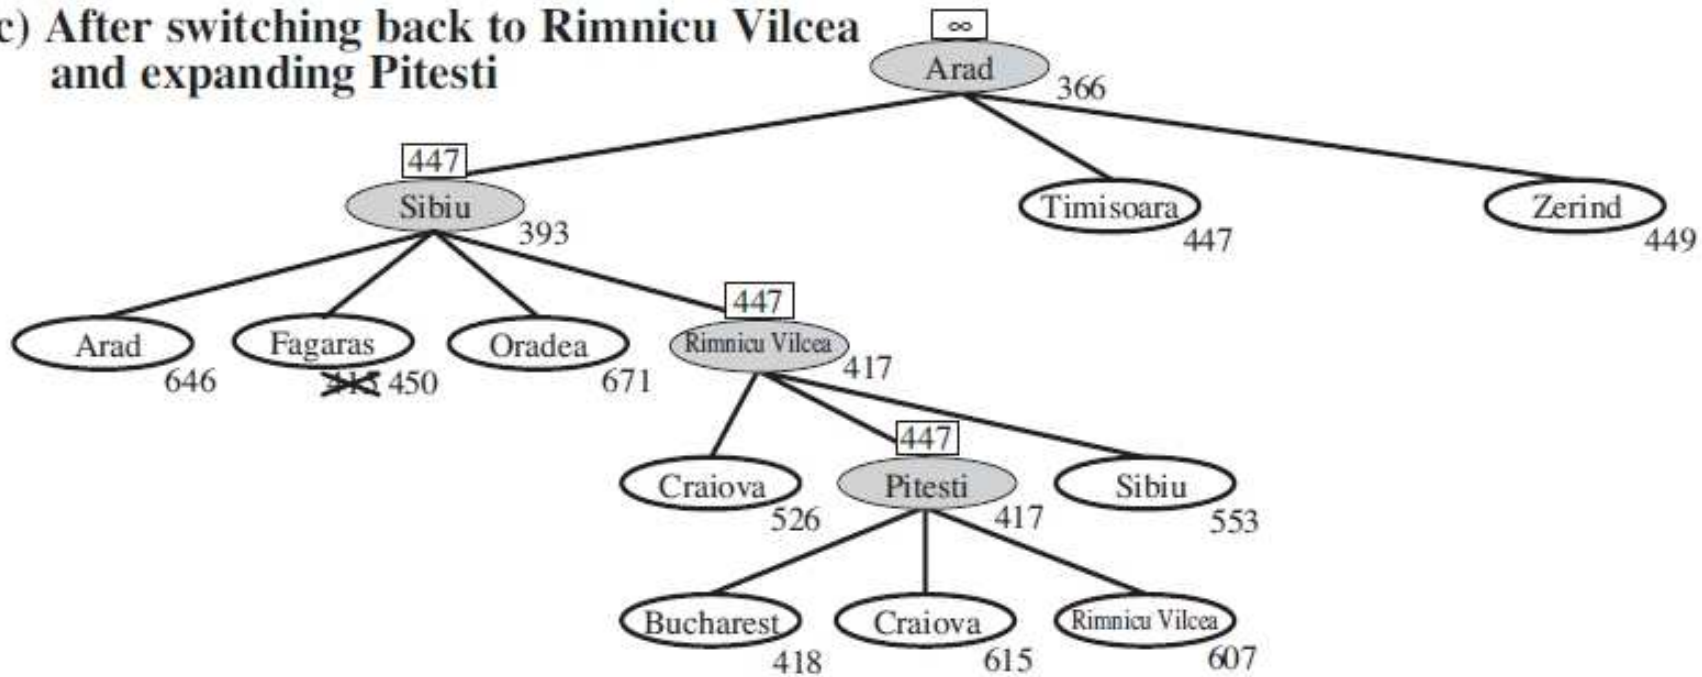
Memory-bounded heuristic search 2

(b) After unwinding back to Sibiu and expanding Fagaras



Memory-bounded heuristic search 3

(c) After switching back to Rimnicu Vilcea and expanding Pitesti



Heuristic functions

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| 1 | 2 | 3 |
| 4 | 5 | 6 |
| 7 | 8 | |

Goal State

h_1 - the number of misplaced tiles

h_2 - the sum of the distances of the tiles from their goal positions

Is it possible for a computer to invent such a heuristic mechanically?

Summary

- ◇ Methods that an agent can use to select actions in environments that are deterministic, observable, static, and completely known. In such cases, the agent can construct sequences of actions that achieve its goals - **search**.
- ◇ Before an agent can start searching for solutions, a goal must be identified and a welldefined problem must be formulated.
- ◇ A problem: 1) initial state, 2) a set of actions, 3) a transition model describing the results of those actions, 4) a goal test function, and 5) a path cost function.
- ◇ Search algorithms: completeness, optimality, time complexity, and space complexity.
- ◇ Uninformed search: breadth-first, uniform cost, depth-first, iterative deepening, bidirectional search
- ◇ Informed (heuristic) search: greedy-best first, A^* , memory bounded